

Vendor-Agnostic RAG + Agentic AI Capability Matrix

Category	Description	Agentic Behaviours	Autonomy Level	Example Use Cases
Self-Learning RAG	Continuously updates retrieval corpus or strategy from feedback	Monitors outcomes, re-indexes data, tunes search parameters	High	Domain knowledge that evolves daily (e.g., regulations, news)
Tool-Calling RAG	Invokes external APIs/tools during retrieval or synthesis	Decides which tool to use, when to enrich data	Medium–High	Fact verification, currency conversion, data transformations
Multi-Turn Agentic RAG	Maintains context across turns and retrieves accordingly	Chooses retrieval scope per turn	Medium	Conversational assistants, research bots
Autonomous Orchestrated RAG	Breaks tasks into subtasks, retrieves for each step	Plans retrieval order, merges results	High	Multi-step research, legal discovery
Context-Adaptive RAG	Adapts retrieval method to query type and intent	Detects query class, changes retrieval backend or filters	Medium	Hybrid search systems, intent-aware customer service
Retrieval-Planning RAG	Creates explicit plan before retrieving	Multi-hop search, branch-and-bound reasoning	Medium–High	Complex investigative search, open-domain Q&A;
Multi-Modal RAG	Retrieves and fuses multiple data modalities	Decides which modality is relevant	Medium–High	Video-text Q&A;, image-aided diagnosis
Verification-Augmented RAG	Validates retrieved facts before output	Self-check, cross-retrieval, voting	Medium	High-stakes knowledge systems (finance, medicine)
Feedback-Driven Iterative RAG	Iteratively refines retrieval until quality threshold met	Self-critique and re-querying	High	Creative writing assistants, technical troubleshooting
Memory-Augmented RAG	Uses persistent memory plus retrieval	Balances memory recall with new retrieval	Medium	Personal AI assistants, longitudinal case tracking
Goal-Oriented RAG	Retrieval focused on achieving a defined objective	Prunes irrelevant search paths	Medium–High	Autonomous research agents, sales intelligence

Extended Attributes

Category	Key Strength	Potential Limitation
Self-Learning RAG	Adaptability to changing information	Risk of model drift if feedback is noisy
Tool-Calling RAG	Extends capabilities beyond text retrieval	Latency and dependency on external systems
Multi-Turn Agentic RAG	Rich contextual understanding	Context window limitations
Autonomous Orchestrated RAG	Handles complex workflows	May overcomplicate simple tasks
Context-Adaptive RAG	Efficient matching to query type	Needs strong intent classification
Retrieval-Planning RAG	Structured and efficient search	Planning errors can block retrieval
Multi-Modal RAG	Supports richer knowledge sources	Complex fusion logic required
Verification-Augmented RAG	High reliability of output	Increased processing time
Feedback-Driven Iterative RAG	Improved final answer quality	Longer response time
Memory-Augmented RAG	Keeps long-term context	Risk of outdated or irrelevant recall
Goal-Oriented RAG	Highly targeted outputs	May miss tangential but valuable insights